

CS2040S

Data Structures and Algorithms

Hashing!
(Part 1)

Plan: this week and next

Three (or Four) Days of Hashing

- Applications
- Basic theory
- Handling collisions
- (Hashing in Java)
- Amortized analysis (doubling/shrinking)
- Sets and Bloom filters

Topic of the Week: Hash Tables

Abstract Data Types

Symbol Table

public interface SymbolTable

| | | |
|----------------------|-------------------------------------|---------------------------------|
| <code>void</code> | <code>insert(Key k, Value v)</code> | <i>insert (k,v) into table</i> |
| <code>Value</code> | <code>search(Key k)</code> | <i>get value paired with k</i> |
| <code>void</code> | <code>delete(Key k)</code> | <i>remove key k (and value)</i> |
| <code>boolean</code> | <code>contains(Key k)</code> | <i>is there a value for k?</i> |
| <code>int</code> | <code>size()</code> | <i>number of (k,v) pairs</i> |

Note: no successor / predecessor queries.

Symbol Table

Examples:

| | |
|---------------|---|
| Dictionary: | key = word value = definition |
| Phone Book | key = name value = phone number |
| Internet DNS | key = website URL value = IP address |
| Java compiler | key = variable name value = type and value |

Implement symbol table with an AVL tree:
(C_I = cost insert, C_S = cost search)

1. $C_I = O(1), C_S = O(1)$
2. $C_I = O(1), C_S = O(\log n)$
3. $C_I = O(1), C_S = O(n)$
- ✓ 4. $C_I = O(\log n), C_S = O(\log n)$
5. $C_I = O(n), C_S = O(\log n)$
6. $C_I = O(n), C_S = O(n)$

Symbol Table

Implement a symbol table with:

- $C_I = O(1)$
- $C_S = O(1)$

Fast, fast, fast....

Dictionaries vs. Symbol Tables

What can you do with a dictionary but not a symbol table?

Dictionaries vs. Symbol Tables

Sorting with a dictionary:

- 1) Insert every item into the dictionary.
- 2) Search for the minimum item.
- 3) Repeat: find successor

Running time to implement sorting:

With an AVL tree/dictionary?

Dictionaries vs. Symbol Tables

Sorting with a dictionary:

- 1) Insert every item into the dictionary.
- 2) Search for the minimum item.
- 3) Repeat: find successor



Running time to implement sorting:

With an AVL tree/dictionary? $O(n \log n)$

With a symbol table?

Dictionaries vs. Symbol Tables

Sorting with a dictionary:

- 1) Insert every item into the dictionary.
- 2) Search for the minimum item.
- 3) Repeat: find successor

Running time to implement sorting:

With an AVL tree/dictionary? $O(n \log n)$

With a symbol table? $O(n^2)$

- No efficient way to find minimum item!
- No ordering of elements.

Sorting (aside)

Isn't $O(1)$ search/insert impossible?

Sorting takes $\Omega(n \log n)$ comparisons.

Sorting (aside)

Isn't $O(1)$ search/insert impossible?

Sorting takes $\Omega(n \log n)$ comparisons.

- How do you sort with a symbol table?
- Only search/insert/delete.

Sorting (aside)

Isn't $O(1)$ search/insert impossible?

Sorting takes $\Omega(n \log n)$ comparisons.

- How do you sort with a symbol table?
- Only search/insert/delete.

(Binary) search takes $\Omega(\log n)$ comparisons.

- Impossible to search in fewer than $\log(n)$ comparisons.
- But a symbol table finds an item in $O(1)$ steps!!
- Conclusion: symbol table is not *comparison-based*.

Building a Symbol Table

Direct Access Tables

Attempt #1: Use a table, indexed by keys.

| | |
|---|-------|
| 0 | null |
| 1 | null |
| 2 | item1 |
| 3 | null |
| 4 | null |
| 5 | item3 |
| 6 | null |
| 7 | null |
| 8 | item2 |
| 9 | null |

Universe $U = \{0..9\}$ of size $m = 10$.

(key, value)

(2, item1)

(8, item2)

(5, item3)


Assume keys are distinct.

Direct Access Tables

Attempt #1: Use a table, indexed by keys.

| | |
|---|-------|
| 0 | null |
| 1 | null |
| 2 | item1 |
| 3 | null |
| 4 | null |
| 5 | item3 |
| 6 | null |
| 7 | null |
| 8 | item2 |
| 9 | null |

Example: insert(4, Seth)




Direct Access Tables

Attempt #1: Use a table, indexed by keys.

| | |
|---|-------|
| 0 | null |
| 1 | null |
| 2 | item1 |
| 3 | null |
| 4 | Seth |
| 5 | item3 |
| 6 | null |
| 7 | null |
| 8 | item2 |
| 9 | null |

Example: insert(4, Seth)



Time: $O(1)$ / insert, $O(1)$ / search

Direct Access Tables

Problems:

- What if keys are not integers?
 - Where do you put the key/value “(hippopotamus, bob)”?
 - Where do you put 3.14159...?

Direct Access Tables

Pythagoras said, “Everything is a number.”



“The School of Athens” by Raphael

Direct Access Tables

Pythagoras said, “Everything is a number.”

- Everything is just a sequence of bits.
- Treat those bits as a number.
- English:
 - 26 letters \Rightarrow 5 bits/letter
 - Longest word = 28 letters (antidisestablishmentarianism?)
 - 28 letters * 5 bits = 140 bits
 - So we can store any English word in a direct-access array of size 2^{140} .

Direct Access Tables

Pythagoras said, “Everything is a number.”

- Everything is just a sequence of bits.
- Treat those bits as a number.
- English:
 - 26 letters \Rightarrow 5 bits/letter
 - Longest word = 28 letters (antidisestablishmentarianism?)
 - 28 letters * 5 bits = 140 bits
 - So we can store any English word in a direct-access array of size 2^{140} . \approx number of atoms in observable universe

Direct Access Tables

Problems:


- What if keys are not integers?
 - Where do you put the key/value “(hippopotamus, bob)”?
 - Where do you put 3.14159...?
 - ➔ Can represent anything as a sequence of bits.
- Too much space
 - If keys are integers, then table-size > 4 billion
 - ➔ Hashing

Hash Functions

Problem:

- Huge universe U of possible keys.
- Smaller number n of actual keys.


e.g., 2^{140}



Hash Functions

Problem:


- Huge universe U of possible keys.
- Smaller number n of actual keys.
- How to map n keys to $m \approx n$ buckets?



Hash Functions

Problem:

- Huge universe U of possible keys.
- Smaller number n of actual keys.
- How to map n keys to $m \approx n$ buckets?




Hash Functions


Define hash function $h : U \rightarrow \{1..m\}$

- Store key k in bucket $h(k)$.
- Time complexity:
 - Time to compute h + Time to access bucket
- Usually: assume hash function has cost 1 to compute.

unless otherwise specified, e.g., long strings.




Hash Functions



| | |
|---|------|
| 0 | null |
| 1 | null |
| 2 | null |
| 3 | null |
| 4 | null |
| 5 | null |
| 6 | null |
| 7 | null |
| 8 | null |
| 9 | null |

Hash Functions


$\text{insert}(k_1, A)$



Hash Functions

$\text{insert}(k_1, A)$

$\text{insert}(k_2, B)$




Hash Functions

$\text{insert}(k_1, A)$

$\text{insert}(k_2, B)$

$\text{insert}(k_3, C)$

Collision!



Hash Functions

Collisions:

- We say that two distinct keys k_1 and k_2 **collide** if:

$$h(k_1) = h(k_2)$$

Can we choose a hash function with no collisions?

1. Yes
2. Sometimes, if we choose carefully
- ✓ 3. No, impossible



Hash Functions

Collisions:

- We say that two distinct keys k_1 and k_2 **collide** if:

$$h(k_1) = h(k_2)$$

- Unavoidable!
 - The table size is smaller than the universe size.
 - The pigeonhole principle says:
 - There must exist two keys that map to the same bucket.
 - Some keys must collide!

Coping with Collision

Idea: choose a new, better hash functions

Coping with Collision

Idea: choose a new, better hash functions

- Hard to find.
- Requires re-copying the table.
- Eventually, there will be another collision.

Coping with Collision

Idea: choose a new, better hash functions

- Hard to find.
- Requires re-copying the table.
- Eventually, there will be another collision.

Idea: chaining (today)

- Put both items in the same bucket!

Idea: open addressing (next week)

- Find another bucket for the new item.

Coping with Collision

Idea: choose a new, better hash functions

- Hard to find.
- Requires re-copying the table.
- Eventually, there will be another collision.

Idea: chaining (today)


- Put both items in the same bucket!

Idea: open addressing (next week)

- Find another bucket for the new item.

Chaining

Each bucket contains a linked list of items.



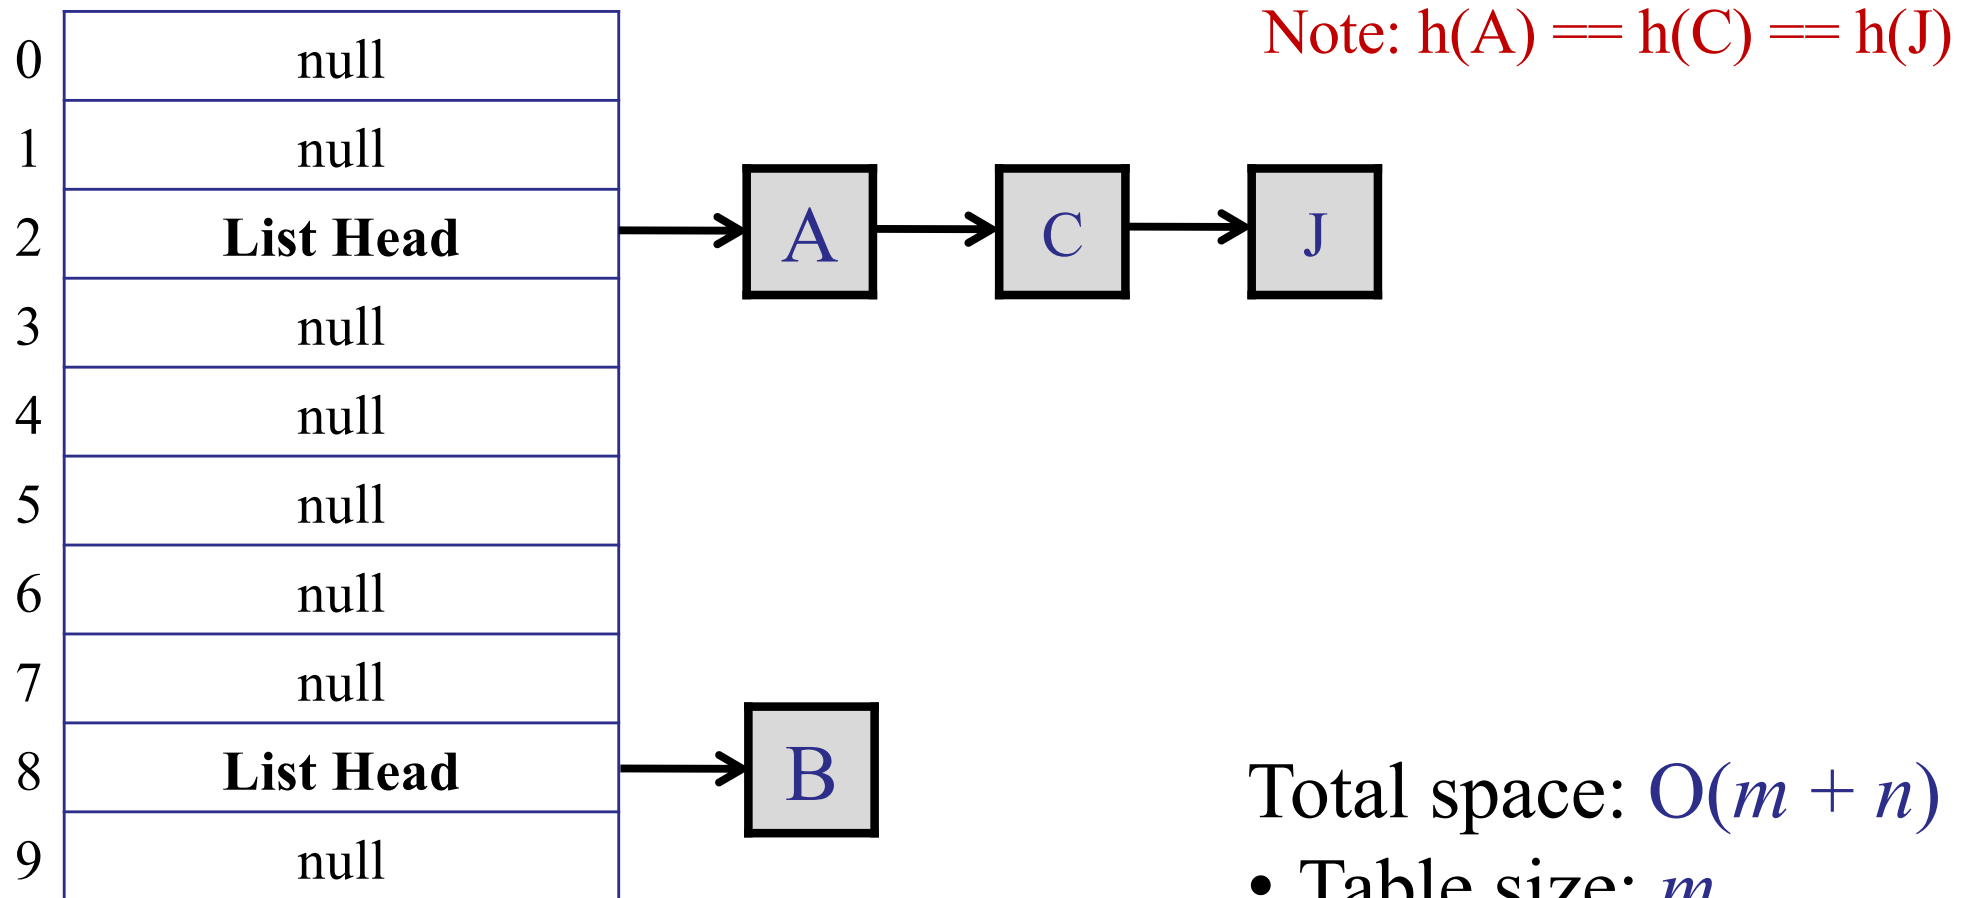
Note: $h(A) == h(C) == h(J)$

Total space:

- Number buckets: m
- Number entries: n

Chaining

Each bucket contains a linked list of items.



Total space: $O(m + n)$

- Table size: m
- Linked list size: n

Hashing with Chaining

Operations:

- insert(key, value)
 - Calculate $h(\text{key})$
 - Lookup $h(\text{key})$ and add (key,value) to the linked list.
- search(key)
 - Calculate $h(\text{key})$
 - Search for (key,value) in the linked list.

What is the worst-case cost of inserting a (key, value)? Assume $\text{cost}(h)$ is cost of computing the hash function.

- ✓ 1. $O(1 + \text{cost}(h))$
- 2. $O(\log n + \text{cost}(h))$
- 3. $O(n + \text{cost}(h))$
- 4. $O(n \text{ cost}(h))$
- 5. $O(n^2)$.

Do we care about duplicates?

→ If so, the cost of insert is higher because we need to search for duplicates.




Hashing with Chaining

Operations:

- `insert(key, value)`
 - Calculate `h(key)`
 - Lookup `h(key)` and add `(key,value)` to the linked list.

Always $O(1)$.



(Note: this allows duplicate keys. Need to specify more precisely the behavior or insert!)

- `search(key)`
 - Calculate `h(key)`
 - Search for `(key,value)` in the linked list.

What is the worst-case cost of searching a (key, value)?

1. $O(1 + \text{cost}(h))$
2. $O(\log n + \text{cost}(h))$
- ✓ 3. $O(n + \text{cost}(h))$
4. $O(n * \text{cost}(h))$
5. We cannot determine it without knowing h .

ARCHIPELAGO

is open

Hashing with Chaining


Operations:

- insert(key, value)
 - Calculate $h(\text{key})$
 - Lookup $h(\text{key})$ and add (key,value) to the linked list.
- search(key) \rightarrow time depends on length of linked list
 - Calculate $h(\text{key})$
 - Search for (key,value) in the linked list.

Hashing with Chaining

Assume all keys hash to the same bucket!

- Search costs $O(n)$
- Oh no!



Let's be optimistic today.


The Simple Uniform Hashing Assumption

- Every key is equally likely to map to every bucket.
- Keys are mapped independently.

Intuition:

- Each key is put in a random bucket.
- Then, as long as there are enough buckets, we won't get too many keys in any one bucket.

Why don't we just insert each key into a random bucket (instead of using a hash function h)?



Why don't we just insert each key into a random bucket (instead of using hash function h)?

1. It would be slow to insert.
2. Computers don't have a real source of randomness.
3. By choosing the keys carefully, a user could force the random choices to create many collisions.
- ✓ 4. Searching would be very slow.
5. None of the above.

Let's be optimistic today.

The Simple Uniform Hashing Assumption

- Assume:
 - n items
 - m buckets
- Define: $\text{load}(\text{hash table}) = n/m$
= average # items / bucket.
- Expected search time = $1 + \text{expected \# items per bucket}$
 - hash function + array access
 - linked list traversal

Probability Theory

Set of outcomes for $X = (e_1, e_2, e_3, \dots, e_k)$:

- $\Pr(e_1) = p_1$
- $\Pr(e_2) = p_2$
- \dots
- $\Pr(e_k) = p_k$

Expected outcome:

$$E[X] = e_1p_1 + e_2p_2 + \dots + e_kp_k$$

Probability Theory

Linearity of Expectation:

- $E[A + B] = E[A] + E[B]$

Example:

- $A = \# \text{ heads in 2 coin flips}$
- $B = \# \text{ heads in 2 coin flips}$
- $A + B = \# \text{ heads in 4 coin flips}$

$$E[A+B] = E[A] + E[B] = 1 + 1 = 2$$

Let's be optimistic today.

The Simple Uniform Hashing Assumption

- Assume:
 - n items
 - m buckets
- Define: $\text{load}(\text{hash table}) = n/m$
= average # items / bucket.
- Expected search time = $1 + \text{expected \# items per bucket}$
 - hash function + array access
 - linked list traversal

A little more probability

Indicator random variables

$$\begin{aligned} X(i, j) &= 1 \quad \text{if item } i \text{ is put in bucket } j \\ &= 0 \quad \text{otherwise} \end{aligned}$$

$$\Pr(X(i, j) == 1) = ?$$

- ✓ 1. $1/m$
- 2. $1/n$
- 3. $1/(m+n)$
- 4. m/n
- 5. n/m
- 6. $\log(n)$

ARCHIPELAGO

is open

Let's be optimistic today.

The Simple Uniform Hashing Assumption

- Every key is equally likely to map to every bucket.
- Keys are mapped independently.

Intuition:

- Each key is put in a random bucket.
- Then, as long as there are enough buckets, we won't get too many keys in any one bucket.

A little probability

Indicator random variables

$X(i, j) = 1$ if item i is put in bucket j
 $= 0$ otherwise

$$\Pr(X(i, j) = 1) = 1/m$$

A little probability

Indicator random variables

$X(i, j) = 1$ if item i is put in bucket j
 $= 0$ otherwise

$$\Pr(X(i, j) = 1) = 1/m$$

$$\mathbf{E}(X(i, j)) = ??$$

A little probability

Indicator random variables

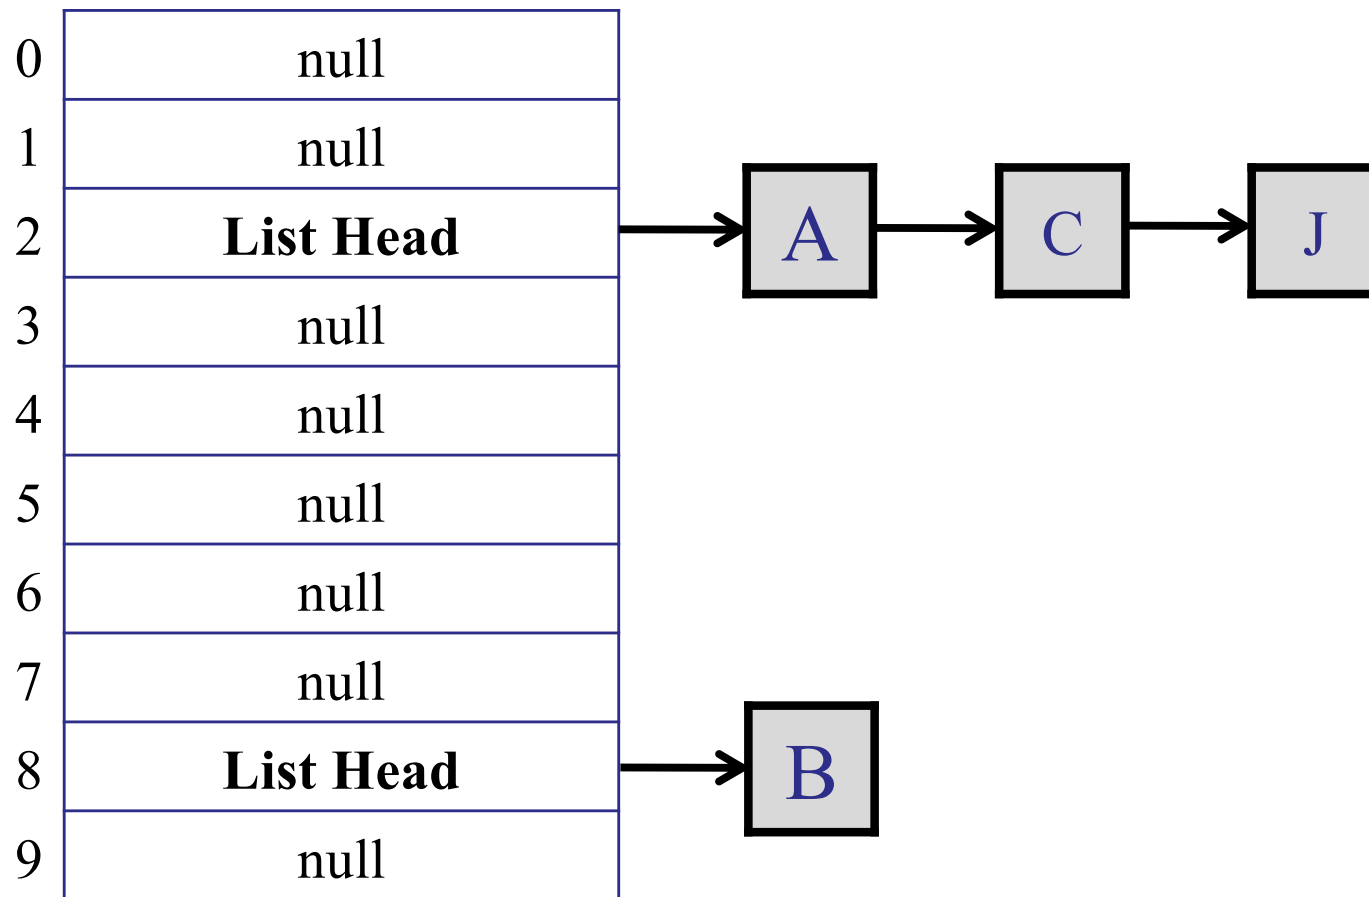
$X(i, j) = 1$ if item i is put in bucket j
 $= 0$ otherwise

$$\mathbf{Pr}(X(i, j) == 1) = 1/m$$

$$\begin{aligned}\mathbf{E}(X(i, j)) &= \mathbf{Pr}(X(i, j) == 1) * 1 + \mathbf{Pr}(X(i, j) == 0) * 0 \\ &= \mathbf{Pr}(X(i, j) == 1) \\ &= 1/m\end{aligned}$$

A little probability

What is the expected number of items in a bucket?



A little probability

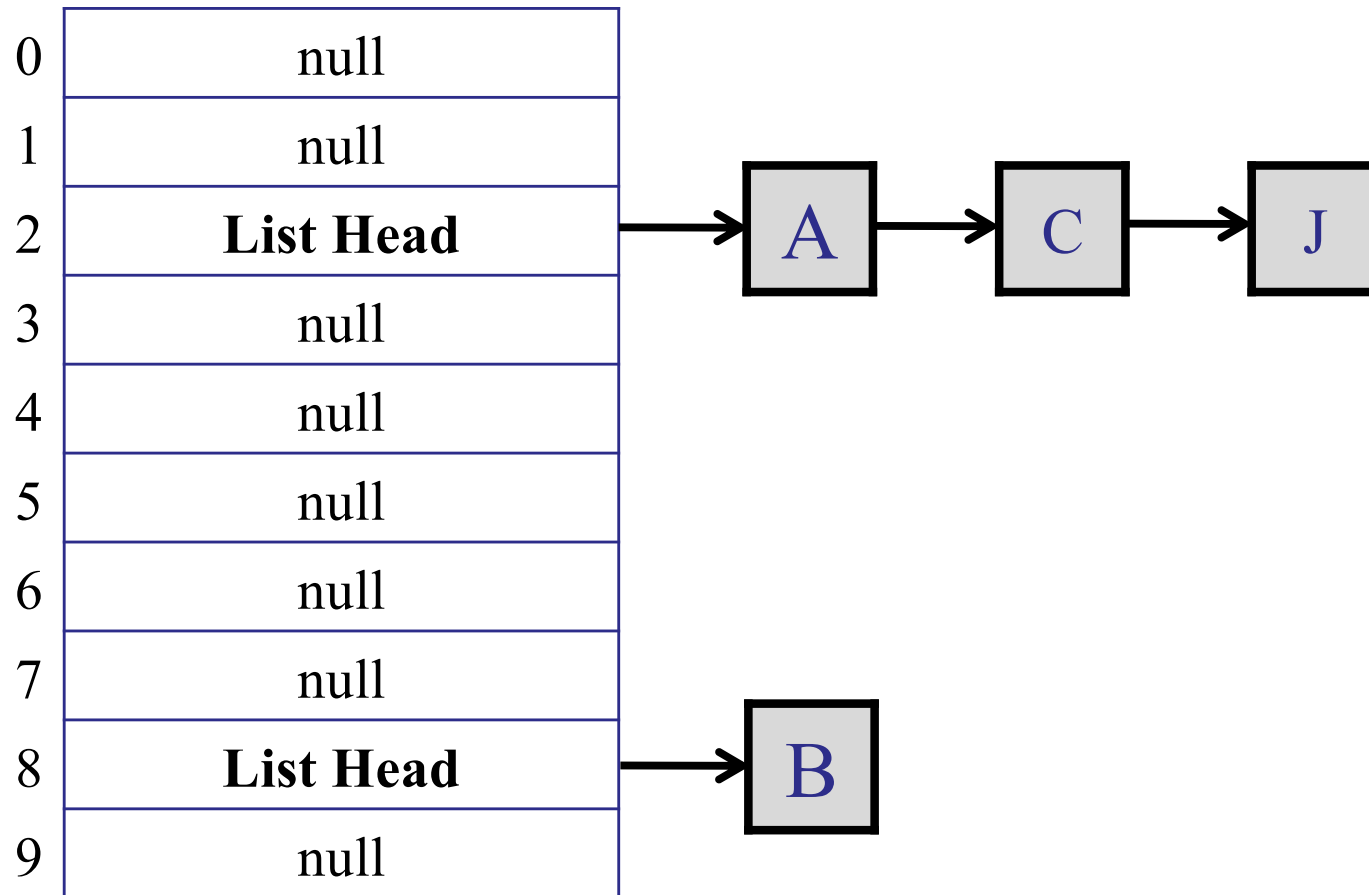
Indicator random variables

$$\begin{aligned} X(i, j) &= 1 \text{ if item } i \text{ is put in bucket } j \\ &= 0 \text{ otherwise} \end{aligned}$$

$$\sum_i X(i, b) = \text{number of items in bucket } b$$

A little probability

Each item contributes `1` to the bucket it is in..



A little probability

Indicator random variables

$$\begin{aligned} X(i, j) &= 1 \text{ if item } i \text{ is put in bucket } j \\ &= 0 \text{ otherwise} \end{aligned}$$

$$\sum_i X(i, b) = \text{number of items in bucket } b$$

A little probability

Calculate expected number of items per bucket:

$$\text{Expected } (\sum_i X(i, b)) =$$

A little probability

Calculate expected number of items per bucket:

$$\mathbf{E}(\sum_i X(i, b)) = \sum_i \mathbf{E} (X(i, b))$$

Linearity of expectation: $E(A + B) = E(A) + E(B)$

A little probability

Calculate expected number of items per bucket:

$$\mathbf{E}(\sum_i X(i, b)) = \sum_i \mathbf{E} (X(i, b))$$

$$= \sum_i 1/m$$

$$= n/m$$

Let's be optimistic today.

The Simple Uniform Hashing Assumption

- Assume:
 - n items
 - m buckets
- Define: $\text{load}(\text{hash table}) = n/m$
 $= \text{average \# items / buckets.}$
- Expected search time = $1 + n/m$
 - hash function + array access
 - linked list traversal

Let's be optimistic today.

The Simple Uniform Hashing Assumption

– Assume:

- n items
- $m = \Omega(n)$ buckets, e.g., $m = 2n$

– Expected search time = $1 + n/m$
= $O(1)$

Hashing with Chaining

Searching:

- Expected search time = $1 + n/m = O(1)$
- Worst-case search time = $O(n)$

Inserting:

- Worst-case insertion time = $O(1)$

Hashing with Chaining

What if you insert n elements in your hash table?

What is the expected *maximum* cost?

Hashing with Chaining

What if you insert n elements in your hash table?

What is the expected *maximum* cost?

– Analogy:

- Throw n balls in $m = n$ bins.
- What is the maximum number of balls in a bin?

Cost: $O(\log n)$

Hashing with Chaining

What if you insert n elements in your hash table?

What is the expected *maximum* cost?

– Analogy:

- Throw n balls in $m = n$ bins.
- What is the maximum number of balls in a bin?

Cost: $\Theta(\log n / \log \log n)$

Hashing: Recap

Problem: coping with large universe of keys

- Number of possible keys is very, very large.
- Direct Access Table takes too much space

Hash functions

- Use hash function to map keys to buckets.
- Sometimes, keys collide (inevitably!)
- Use linked list to store multiple keys in one bucket.

Analyze performance with simple uniform hashing.

- Expected number of keys / bucket is $O(n/m) = O(1)$.